


ANT 6973: DATA VISUALIZATION AND EXPLORATION

DATA MANIPULATION, PART 2

SINGLE TABLE VERBS




Extract variables with `select()`



Extract cases with `filter()`



Arrange cases with `arrange()`



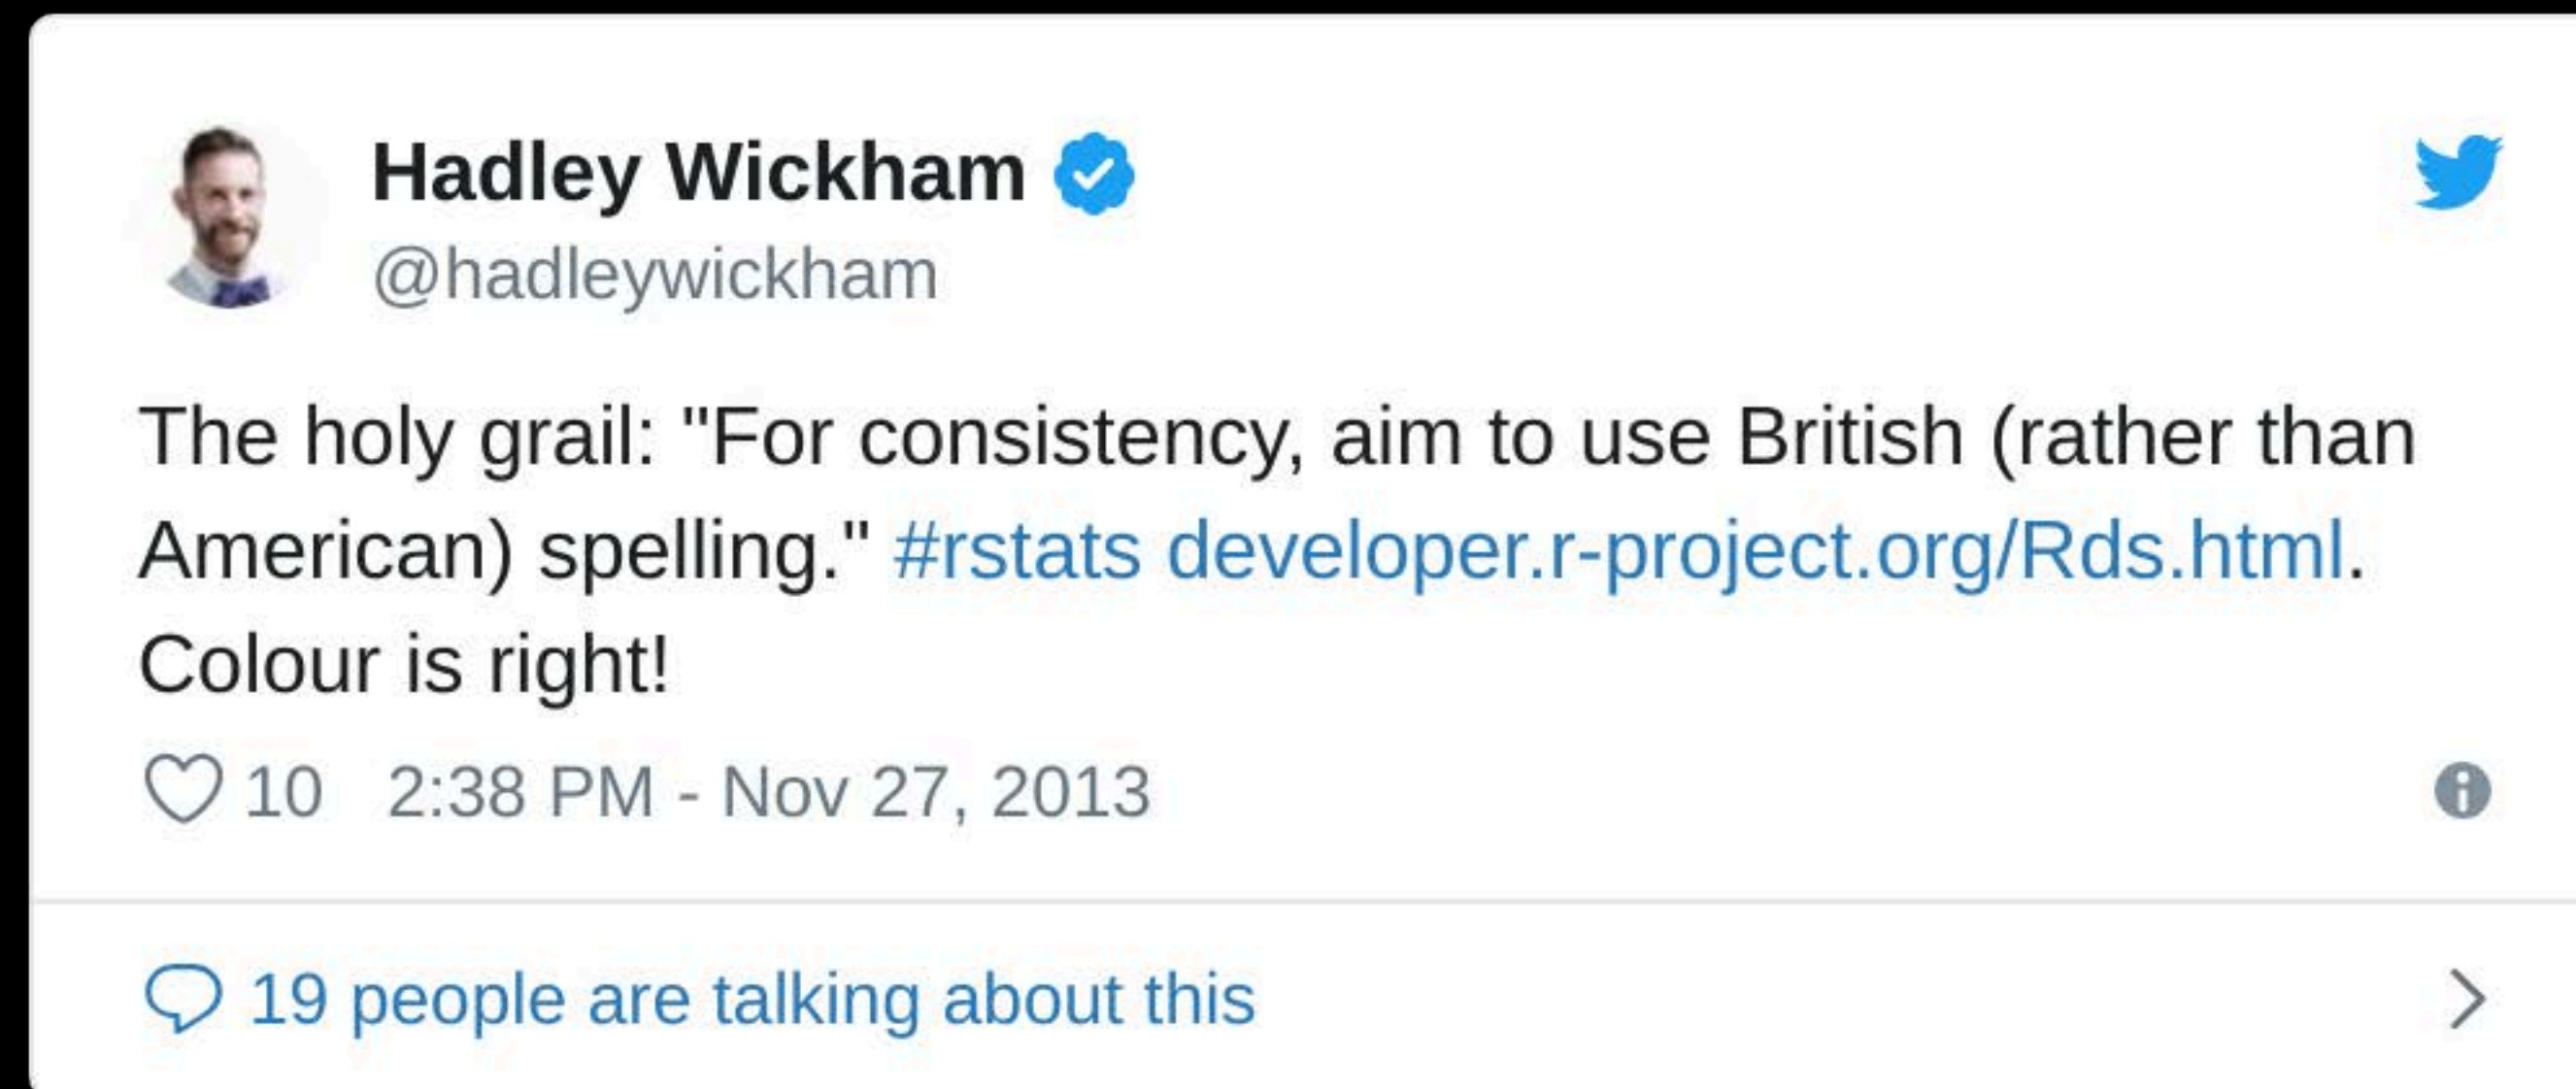
Make new variables with `mutate()`



Make tables of summaries with `summarise()`
along with `group_by()`

summarise()

A NOTE ON SPELLING...





But both the British and American spellings work:

`summarize() = summarise()`

`color() = colour()`

SUMMARISE()

Compute table of summaries.



A *summary function* returns a *single value* that summarizes many values in a column



SUMMARISE()

Compute table of summaries.




```
summarise(.data, ...)
```

data frame to
summarize

Name-value pairs of
summary functions

SUMMARISE()

Compute table of summaries.



```
summarise(babynames, max_prop = max(prop))
```

data frame to
summarize


Name of new
variable

Summary
function

Column of
data to
summarize


SUMMARISE()

Compute table of summaries.



```
summarise(babynames, max_prop = max(prop))
```


year	sex	name	n	prop
1880	F	Mary	7065	0.07238
1880	F	Anna	2604	0.02667
1880	F	Emma	2003	0.02052
1880	F	Elizabet	1939	0.01986
1880	F	Minnie	1746	0.01788
1880	F	Margare	1578	0.01616



max_prop
0.0815

SUMMARISE()

Compute table of summaries.



```
babynames %>%  
  summarise(max_n = max(n),  
            min_n = min(n))
```

Multiple summaries
separated by commas

year	sex	name	n	prop
1880	F	Mary	7065	0.07238
1880	F	Anna	2604	0.02667
1880	F	Emma	2003	0.02052
1880	F	Elizabet	1939	0.01986
1880	F	Minnie	1746	0.01788
1880	F	Margare	1578	0.01616



max_n	min_n
99686	5

USEFUL SUMMARY FUNCTIONS

- **Center:** `mean()`, `median()`
- **Spread:** `sd()`, `var()`
- **Range:** `min()`, `max()`, `quantile()`
- **Logical:** `any()`, `all()`
- **Position:** `first()`, `last()`, `nth()`
- **Count:** `n()`, `n_distinct()`

USEFUL SUMMARY FUNCTIONS

- **Center:** `mean()`, `median()`
- **Spread:** `sd()`, `var()`
- **Range:** `min()`, `max()`, `quantile()`
- **Logical:** `any()`, `all()`
- **Position:** `first()`, `last()`, `nth()`
- **Count:** `n()`, `n_distinct()`

Base R functions

USEFUL SUMMARY FUNCTIONS

- **Center:** `mean()`, `median()`
- **Spread:** `sd()`, `var()`
- **Range:** `min()`, `max()`, `quantile()`
- **Logical:** `any()`, `all()`
- **Position:** `first()`, `last()`, `nth()`
- **Count:** `n()`, `n_distinct()`

dplyr functions

ACTIVITY 1

- Use `summarise()` to compute three statistics about babynames:
 - The smallest (minimum) year in the dataset
 - The largest (maximum) year in the dataset
 - The total number of children represented in the data

```
babynames %>%  
  summarise(first_yr = min(year),  
            last_yr = max(year),  
            total_n = sum(n))
```

```
first_yr last_yr total_n  
  <dbl>   <dbl>    <int>  
1    1880    2017 348120517
```

ACTIVITY 2

- Extract the rows where name is "Khaleesi". Then use `summarise()` to find:
 - The total number of children named Khaleesi
 - The first year Khaleesi appeared in the data

```
babynames %>%  
  filter(name == "Khaleesi") %>%  
  summarise(total = sum(n),  
            first_year = min(year))  
  
total first_year  
<int> <dbl>  
1    1964    2011
```



USEFUL SUMMARY FUNCTIONS

- **Center:** `mean()`, `median()`
- **Spread:** `sd()`, `var()`
- **Range:** `min()`, `max()`, `quantile()`
- **Logical:** `any()`, `all()`
- **Position:** `first()`, `last()`, `nth()`
- **Count:** `n()`, `n_distinct()`

dplyr functions

n()

The number of rows in a dataset/group



```
babynames %>%  
  summarise(n_rows = n())
```


year	sex	name	n	prop
1880	F	Mary	7065	0.07238
1880	F	Anna	2604	0.02667
1880	F	Emma	2003	0.02052
1880	F	Elizabet	1939	0.01986
1880	F	Minnie	1746	0.01788
1880	F	Margare	1578	0.01616



n_rows
1924665

n_distinct()

The number of distinct values in a column/group.



```
babynames %>%  
  summarise(n_rows = n(),  
            n_names = n_distinct(name))
```

year	sex	name	n	prop
1880	F	Mary	7065	0.07238
1880	F	Anna	2604	0.02667
1880	F	Emma	2003	0.02052
1880	F	Elizabet	1939	0.01986
1880	F	Minnie	1746	0.01788
1880	F	Margare	1578	0.01616



n_rows	n_names
1924665	97310

Grouping cases

GROUP_BY()

- `group_by()` changes the unit of analysis to *groups* in the data
- Any `dplyr` verbs used on a *grouped tibble* will be applied "by group"
- Especially useful when paired with `summarise()`
- Use `ungroup()` to remove the grouping

GROUP_BY()

Groups cases by common values of one or more columns.

```
babynames %>%  
  group_by(sex)
```

Source: local data frame [1,825,433 x 5]

Groups: sex [2]

	year	sex	name	n	prop
	<dbl>	<chr>	<chr>	<int>	<dbl>
1	1880	F	Mary	7065	0.07238359

GROUP_BY()

Groups cases by common values of one or more columns.

```
babynames %>%  
  group_by(sex) %>%  
  summarise(total = sum(n))
```

Grouping variable

sex	total
F	172371079
M	175749438

New summary variable

Note that all other columns in original data are not in summary

PRACTICE DATA

```
resave.Rmd* x
~/OneDrive - University of Texas at San Antonio/Teaching/Data Visualization/activities/ant6973-activities - RStudio Source Editor


1 ---
2 title: "Tidy Data"
3 output: html_document
4 editor_options:
5   chunk_output_type: console
6 ---
7
8 ```{r setup, include=FALSE}
9 knitr::opts_chunk$set(echo = TRUE)
10
11 library("gapminder")
12 library("tidyverse")
13 library("knitr")
14 ```
15
16
17 ```{r}
18 cases <- tibble(country = c("FR", "DE", "US"),
19                 `2011` = c(7000, 5800, 15000),
20                 `2012` = c(6900, 6000, 14000),
21                 `2013` = c(7000, 6200, 13000))
22 ```
23
24
25 ```{r}
26 pollution <- tibble(city = c("New York", "New York", "New York", "New York", "New York"),
27                     size = c("large", "small", "large", "small", "large"),
28                     amount = c(23, 14, 22, 16, 12))
29 ```
30
```

```
pollution <- tibble(city = ...,
                     size = ...,
                     amount = ...)
```




```
pollution %>%  
  summarise(mean = mean(amount), sum = sum(amount), n = n())
```

city	particle size	amount (µg/m³)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56




mean	sum	n
42	252	6

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14



mean	sum	n
18.5	37	2

London	large	22
London	small	16



19.0	38	2
------	----	---

Beijing	large	121
Beijing	small	56



88.5	177	2
------	-----	---

`group_by()` + `summarise()`

city	particle size	amount (µg/m³)
New York	large	23
New York	small	14

London	large	22
London	small	16

Beijing	large	121
Beijing	small	56

mean	sum	n
18.5	37	2
19.0	38	2
88.5	177	2

`group_by()` + `summarise()`

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14

London	large	22
London	small	16

Beijing	large	121
Beijing	small	56

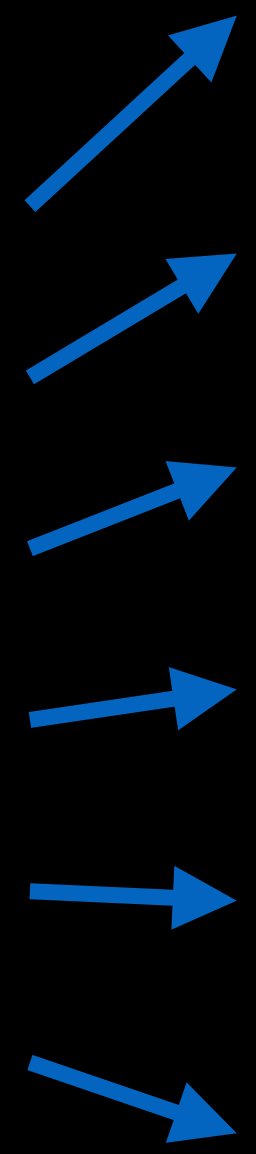
city	mean	sum	n
New York	18.5	37	2
London	19.0	38	2
Beijing	88.5	177	2

```
pollution %>%
```

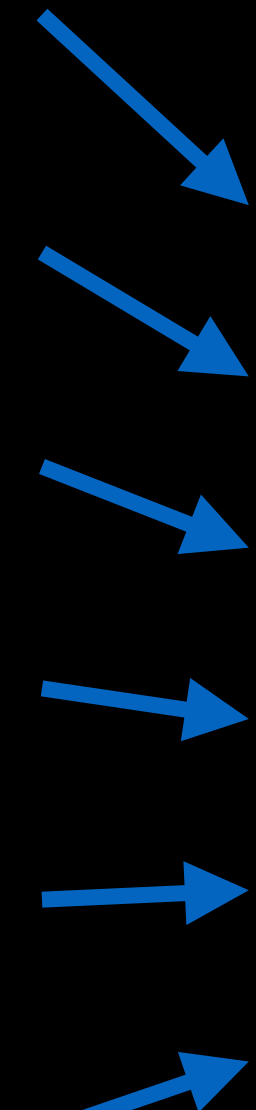
```
  group_by(city) %>%
```

```
  summarise(mean = mean(amount), sum = sum(amount), n = n())
```

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



city	particle size	mean	sum	n
New York	large	23	23	1
New York	small	14	14	1
London	large	22	22	1
London	small	16	16	1
Beijing	large	121	121	1
Beijing	small	56	56	1

```
pollution %>%
```

```
  group_by(city, size) %>%
```

```
  summarise(mean = mean(amount), sum = sum(amount), n = n())
```

ACTIVITY 3

- Use `group_by()` and `summarise()` to calculate the number of male and female babies born in each year.

```
babynames %>%  
  group_by(year, sex) %>%  
  summarise(n_babies = sum(n))
```

```
# Groups:   year [138]  
   year sex  n_babies  
   <dbl> <chr>    <int>  
1  1880 F      90993  
2  1880 M     110491  
3  1881 F      91953  
4  1881 M     100743  
5  1882 F     107847  
6  1882 M     113686  
7  1883 F     112319  
8  1883 M     104627  
9  1884 F     129020  
10 1884 M     114442  
# ... with 266 more rows
```

Note that each call to `summarise()` removes a layer of grouping.

ACTIVITY 4

- On the storms data set, calculate the maximum wind speed and minimum pressure for each **hurricane** in each year, and arrange the summary in descending order of wind speed (hint: filter first).

name	year	month	day	hour	lat	long	status	category	wind	pressure	ts_diameter	hu_diameter
Amy	1975	6	27	0	27.5	-79.0	tropical depression	-1	25	1013	NA	NA
Amy	1975	6	27	6	28.5	-79.0	tropical depression	-1	25	1013	NA	NA
Amy	1975	6	27	12	29.5	-79.0	tropical depression	-1	25	1013	NA	NA
Amy	1975	6	27	18	30.5	-79.0	tropical depression	-1	25	1013	NA	NA
Amy	1975	6	28	0	31.5	-78.8	tropical depression	-1	25	1012	NA	NA
Amy	1975	6	28	6	32.4	-78.7	tropical depression	-1	25	1012	NA	NA


```
storms %>%
  filter(status == "hurricane") %>%
  group_by(year, name) %>%
  summarise(wind_max = max(wind),
            pressure_min = min(pressure)) %>%
  arrange(desc(wind_max))
```

```
# Groups:   year [41]
```

	year	name	wind_max	pressure_min
	<dbl>	<chr>	<dbl>	<dbl>
1	1988	Gilbert	160	888
2	2005	Wilma	160	882
3	1998	Mitch	155	905
4	2005	Rita	155	895
5	1977	Anita	150	926
6	1979	David	150	924
7	1992	Andrew	150	922
8	2005	Katrina	150	902

```
# ... with 200 more rows
```

ACTIVITY 5

- Building on the previous code, calculate the average hurricane wind_max for each year. Which year had the most intense hurricanes, on average?

```
storms %>%  
  filter(status == "hurricane") %>%  
  group_by(year, name) %>%  
  summarise(wind_max = max(wind)) %>%  
  summarise(avg_wind_max = mean(wind_max)) %>%  
  arrange(desc(avg_wind_max))
```

	year	avg_wind_max
	<dbl>	<dbl>
1	1999	133.
2	1988	117.
3	1992	113.
4	2008	108.
5	2004	104.
6	2005	103.
7	2009	103.
8	2002	98.8
#	...	with 33 more rows

Second call to `summarise()` uses the year grouping only.

You might want to make the groupings explicit for readability.

```

storms %>%
  filter(status == "hurricane") %>%
  group_by(year, name) %>%
  summarise(wind_max = max(wind)) %>%
  ungroup() %>%
  group_by(year) %>%
  summarise(avg_wind_max = mean(wind_max))
  arrange(desc(avg_wind_max))

```

	year	avg_wind_max
	<dbl>	<dbl>
1	1999	133.
2	1988	117.
3	1992	113.
4	2008	108.
7	2009	103.
8	2002	98.8

```
# ...
```

Second call to `summarise()` uses the year grouping only.

You might want to make the groupings explicit for readability.

THINGS TO WATCH OUT FOR

- Many summary functions will return NA if there are any missing values
- Fortunately, many summary functions have an `na.rm = TRUE` argument to avoid this problem.

THINGS TO WATCH OUT FOR

- Once grouped, the tibble will remain grouped. Careful with carrying out further summaries!
- It's good practice to always use `ungroup()` after finishing your grouped operations.

ACTIVITY 6

- What is the average diameter of the area experiencing hurricane strength winds (`hu_diameter`) for each category of hurricane (`category`)?

```
storms %>%
```

```
  group_by(category) %>%
```

```
  summarise(mean_diameter = mean(hu_diameter))
```

```
category mean_diameter
```

```
<ord>
```

```
<dbl>
```

```
1 -1 NA
```

```
2 0 NA
```

```
3 1 NA
```

```
4 2 NA
```

```
5 3 NA
```

```
6 4 NA
```

```
7 5 NA
```

Many of these values
are missing

This causes the means
to be NA


```
storms %>%  
  group_by(category) %>%  
  summarise(mean_diameter = mean(hu_diameter, na.rm = TRUE)))
```

	category	mean_diameter
	<ord>	<dbl>
1	-1	0
2	0	0
3	1	57.3
4	2	78.8
5	3	91.4
6	4	102.
7	5	120.

Remove missing values
before calculating
means

SUMMARISE()'S VARIANTS

- Similarly to the `filter()`, `select()` and `mutate()` functions, `summarise()` comes with three additional variants for doing things to multiple columns in one go:
 - `summarise_all()` will summarise all columns with the summary function(s).
 - `summarise_if()` requires a function that returns a TRUE/FALSE. If that is true, the summary function(s) will be applied.
 - `summarise_at()` requires you to specify columns on which to apply the summary function(s) inside a `vars()` argument.

msleep (lots of missing values)

name	genus	vore	order	conservation	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
Cheetah	Acinonyx	carni	Carnivora	lc	12.1	NA	NA	11.9	NA	50.000
Owl monkey	Aotus	omni	Primates	NA	17.0	1.8	NA	7.0	0.01550	0.480
Mountain beaver	Aplodontia	herbi	Rodentia	nt	14.4	2.4	NA	9.6	NA	1.350
Greater short-tailed shrew	Blarina	omni	Soricomorpha	lc	14.9	2.3	0.1333333	9.1	0.00029	0.019
Cow	Bos	herbi	Artiodactyla	domesticated	4.0	0.7	0.6666667	20.0	0.42300	600.000
Three-toed sloth	Bradypus	herbi	Pilosa	NA	14.4	2.2	0.7666667	9.6	NA	3.850
Northern fur seal	Callorhinus	carni	Carnivora	vu	8.7	1.4	0.3833333	15.3	NA	20.490
Vesper mouse	Calomys	NA	Rodentia	NA	7.0	NA	NA	17.0	NA	0.045
Dog	Canis	carni	Carnivora	domesticated	10.1	2.9	0.3333333	13.9	0.07000	14.000
Roe deer	Capreolus	herbi	Artiodactyla	lc	3.0	NA	NA	21.0	0.09820	14.800

ACTIVITY 7A

- For each vore, calculate the average of *all* columns in the msleep data set.
 - Hint: use `na.rm = TRUE`

```
msleep %>%
```

```
group_by(vore) %>%
```

```
summarise_all(mean, na.rm = TRUE)
```

	vore	name	genus	order	conservation	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	NA	NA	NA	NA	NA	10.2	1.88	0.183	13.8	0.00763	0.858
2	carni	NA	NA	NA	NA	10.4	2.29	0.373	13.6	0.0793	90.8
3	herbi	NA	NA	NA	NA	9.51	1.37	0.418	14.5	0.622	367.
4	insecti	NA	NA	NA	NA	14.9	3.52	0.161	9.06	0.0216	12.9
5	omni	NA	NA	NA	NA	10.9	1.96	0.592	13.1	0.146	12.7

**Note that mean()
doesn't make sense for
the text variables.**

ACTIVITY 7B

- For each vore, calculate the average of *only the numeric columns* in the msleep data set.
- Hint: use `na.rm = TRUE`

```
msleep %>%
```

```
  group_by(vore) %>%
```

```
  summarise_if(is.numeric, mean, na.rm = TRUE)
```

	vore	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	NA	10.2	1.88	0.183	13.8	0.00763	0.858
2	carni	10.4	2.29	0.373	13.6	0.0793	90.8
3	herbi	9.51	1.37	0.418	14.5	0.622	367.
4	insecti	14.9	3.52	0.161	9.06	0.0216	12.9
5	omni	10.9	1.96	0.592	13.1	0.146	12.7

Grouping variable isn't
numeric, but it still
appears in the
summary.

ACTIVITY 7C

- For each vore, calculate the average of *only the columns that start with "sleep"* in the msleep data set.
- Hint: use `na.rm = TRUE`


```
msleep %>%
```

```
  group_by(vore) %>%
```

```
  summarise_at(vars(starts_with("sleep")), mean, na.rm = TRUE)
```

```
vore      sleep_total sleep_rem sleep_cycle
```

```
<chr>
```

```
1 NA
```

```
2 carni
```

```
3 herbi
```

```
4 insecti
```

```
5 omni
```

The `_at()` functions expect a
`select()`-style column
specification within `vars()`

```
14.9
```

```
3.52
```

```
0.161
```

```
10.9
```

```
1.96
```

```
0.592
```

Equivalent to...

```
msleep %>%
```


```
  group_by(vore) %>%
```

```
  summarise(sleep_total = mean(sleep_total, na.rm = TRUE),  
            sleep_rem = mean(sleep_rem, na.rm = TRUE),  
            sleep_cycle = mean(sleep_cycle, na.rm = TRUE))
```

	vore	sleep_total	sleep_rem	sleep_cycle
	<chr>	<dbl>	<dbl>	<dbl>
1	NA	10.2	1.88	0.183
2	carni	10.4	2.29	0.373
3	herbi	9.51	1.37	0.418
4	insecti	14.9	3.52	0.161
5	omni	10.9	1.96	0.592

OTHER GROUPED OPERATIONS

- `group_by()` can also be used with `mutate()` and `filter()` to do some interesting things.
- Reminder: `mutate()` creates a new variable of the same length as the original data.



ACTIVITY 8

- In babynames, try to recreate the prop column using a grouped mutate (call it "new_prop"). Specifically, divide each row's n by the total number of n for that sex and year.
- Why are the values slightly different?

```
babynames %>%  
  group_by(sex, year) %>%  
  mutate(grp_sum = sum(n),  
         new_prop = n / grp_sum)
```

Current *group's*
total n

Current row's
new_prop

Groups: sex, year [276]

	year	sex	name	n	prop	grp_sum	new_prop
	<dbl>	<chr>	<chr>	<int>	<dbl>	<int>	<dbl>
1	1880	F	Mary	7065	0.0724	90993	0.0776
2	1880	F	Anna	2604	0.0267	90993	0.0286
3	1880	F	Emma	2003	0.0205	90993	0.0220
4	1880	F	Elizabeth	1939	0.0199	90993	0.0213
5	1880	F	Minnie	1746	0.0179	90993	0.0192
6	1880	F	Margaret	1578	0.0162	90993	0.0173
7	1880	F	Ida	1472	0.0151	90993	0.0162
8	1880	F	Alice	1414	0.0145	90993	0.0155
9	1880	F	Bertha	1320	0.0135	90993	0.0145
10	1880	F	Sarah	1288	0.0132	90993	0.0142

... with 1,924,655 more rows

Note that the
result has as
many rows as the
original data (it's
not a summary)

`babynames %>%`

`group_by(sex, year) %>%`

`mutate(new_prop = n / sum(n))`

More concisely...

Group denominator
calculated for each for
each row

Groups: sex, year [276]

	year	sex	name	n	prop	new_prop
	<dbl>	<chr>	<chr>	<int>	<dbl>	<dbl>
1	1880	F	Mary	7065	0.0724	0.0776
2	1880	F	Anna	2604	0.0267	0.0286
3	1880	F	Emma	2003	0.0205	0.0220
4	1880	F	Elizabeth	1939	0.0199	0.0213
5	1880	F	Minnie	1746	0.0179	0.0192
6	1880	F	Margaret	1578	0.0162	0.0173
7	1880	F	Ida	1472	0.0151	0.0162
8	1880	F	Alice	1414	0.0145	0.0155
9	1880	F	Bertha	1320	0.0135	0.0145
10	1880	F	Sarah	1288	0.0132	0.0142

... with 1,924,655 more rows

MORE ON GROUPED MUTATES

- Grouped mutates are useful for calculating deviations.

ACTIVITY 9

- Using msleep, determine how much each species' log body weight differs from the average log body weight for its order.

name	genus	vore	order	conservation	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
Cheetah	Acinonyx	carni	Carnivora	lc	12.1	NA	NA	11.9	NA	50.000
Owl monkey	Aotus	omni	Primates	NA	17.0	1.8	NA	7.0	0.01550	0.480
Mountain beaver	Aplodontia	herbi	Rodentia	nt	14.4	2.4	NA	9.6	NA	1.350
Greater short-tailed shrew	Blarina	omni	Soricomorpha	lc	14.9	2.3	0.1333333	9.1	0.00029	0.019
Cow	Bos	herbi	Artiodactyla	domesticated	4.0	0.7	0.6666667	20.0	0.42300	600.000
Three-toed sloth	Bradypus	herbi	Pilosa	NA	14.4	2.2	0.7666667	9.6	NA	3.850
Northern fur seal	Callorhinus	carni	Carnivora	vu	8.7	1.4	0.3833333	15.3	NA	20.490

msleep %>%

group_by(order) %>%

mutate(log_bodywt = log(bodywt),
order_mean = mean(log_bodywt, na.rm = TRUE),
bodywt_dev = log_bodywt - order_mean)

Each row's log
bodywt

Current order's
mean bodywt

Current row's
deviation from mean
of current order

Groups: order [19]


	name	genus	vore	order	conservation	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt	log_bodywt	order_mean
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	Cheetah	Acinon...	carni	Carniv...	lc	12.1	NA	NA	11.9	NA	50	3.91	3.91
2	Owl monkey	Aotus	omni	Primat...	NA	17	1.8	NA	7	0.0155	0.48	-0.734	1.17
3	Mountain beav...	Aplodo...	herbi	Rodent...	nt	14.4	2.4	NA	9.6	NA	1.35	0.300	-1.98
4	Greater short...	Blarina	omni	Sorico...	lc	14.9	2.3	0.133	9.1	0.00029	0.019	-3.96	-3.54
5	Cow	Bos	herbi	Artiod...	domesticated	4	0.7	0.667	20	0.423	600	6.40	4.65
6	Three-toed sl...	Bradyp...	herbi	Pilosa	NA	14.4	2.2	0.767	9.6	NA	3.85	1.35	1.35
7	Northern fur ...	Callor...	carni	Carniv...	vu	8.7	1.4	0.383	15.3	NA	20.5	3.02	3.15
8	Vesper mouse	Calomys	NA	Rodent...	NA	7	NA	NA	17	NA	0.045	-3.10	-1.98
9	Dog	Canis	carni	Carniv...	domesticated	10.1	2.9	0.333	13.9	0.07	14	2.64	3.15
10	Roe deer	Capreo...	herbi	Artiod...	lc	3	NA	NA	21	0.0982	14.8	2.69	4.65

ACTIVITY 10A

- When building statistical models, it's often useful to standardize numeric variables by converting to z-scores:
 - $z = (\text{observed value} - \text{mean}) / \text{standard deviation}$
- Using `msleep`, use `mutate()` to convert `log(bodywt)` to a z-score for each row in the data set.
- Plot the data using a violin plot and `stat_summary()`

```
ggplot(msleep, aes(x = vore, y = z_log_bodywt)) +  
  geom_violin() +  
  stat_summary(fun.data = "mean_se")
```

```
msleep %>%  
  mutate(log_bodywt = log(bodywt),  
         z_log_bodywt = (log_bodywt - mean(log_bodywt, na.rm = TRUE)) /  
                        sd(log_bodywt, na.rm = TRUE)) %>%  
  ggplot(aes(x = vore, y = z_log_bodywt)) +  
  geom_violin() +  
  stat_summary(fun.data = "mean_se")
```




ACTIVITY 10B

- Now do the same, but standardize *within groups* by using `group_by(vore)`.

```
ggplot(msleep, aes(x = vore, y = z_log_bodywt)) +  
  geom_violin() +  
  stat_summary(fun.data = "mean_se")
```


```
msleep %>%  
  mutate(log_bodywt = log(bodywt)) %>%  
  group_by(vore) %>%  
  mutate(z_log_bodywt = (log_bodywt - mean(log_bodywt, na.rm = TRUE)) /  
                                sd(log_bodywt, na.rm = TRUE)) %>%  
  ggplot(aes(x = vore, y = z_log_bodywt)) +  
  geom_violin() +  
  stat_summary(fun.data = "mean_se")
```



SINGLE TABLE VERBS




Extract variables with `select()`



Extract cases with `filter()`



Arrange cases with `arrange()`



Make new variables with `mutate()`



Make tables of summaries with `summarise()`
along with `group_by()`

ACKNOWLEDGEMENTS

- Some ideas, examples, and figures from RStudio webinars, which are licensed CC by SA.